

II - Linear and Logistic Regression

Formalism & Inference function

Loss function

Algorithms

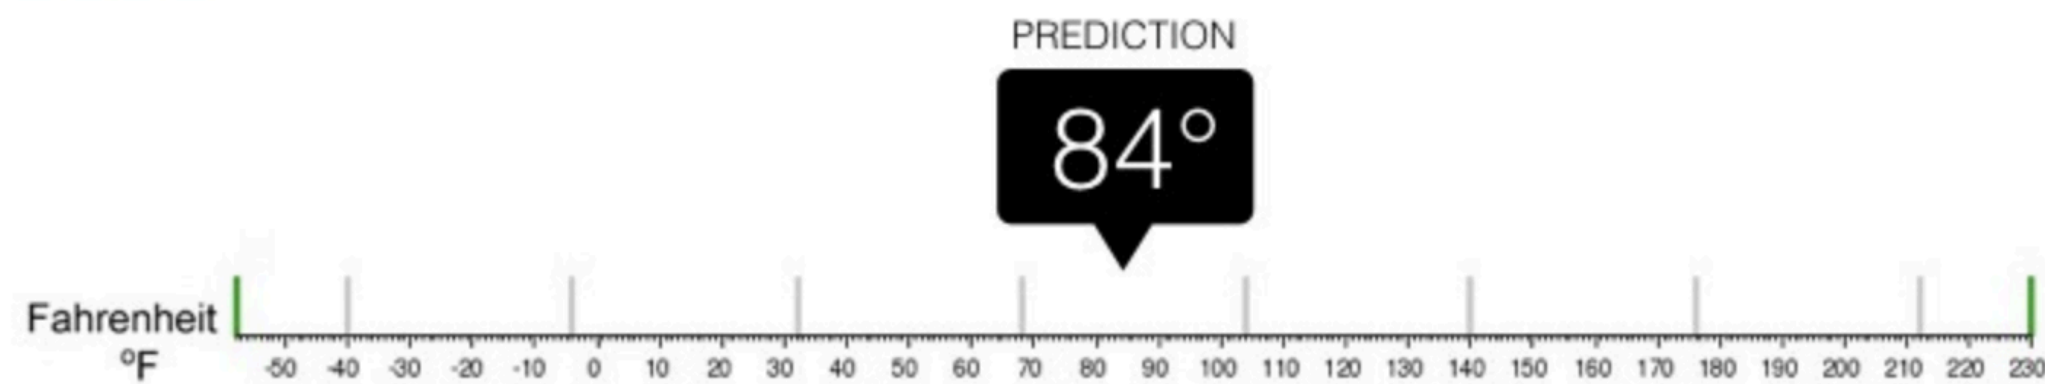
youssef.ait-el-mahjoub@efrei.fr

Regression Vs Classification




Regression

What is the temperature going to be tomorrow?





Classification

Will it be Cold or Hot tomorrow?



Source: [Pinterest](#)



Linear Regression

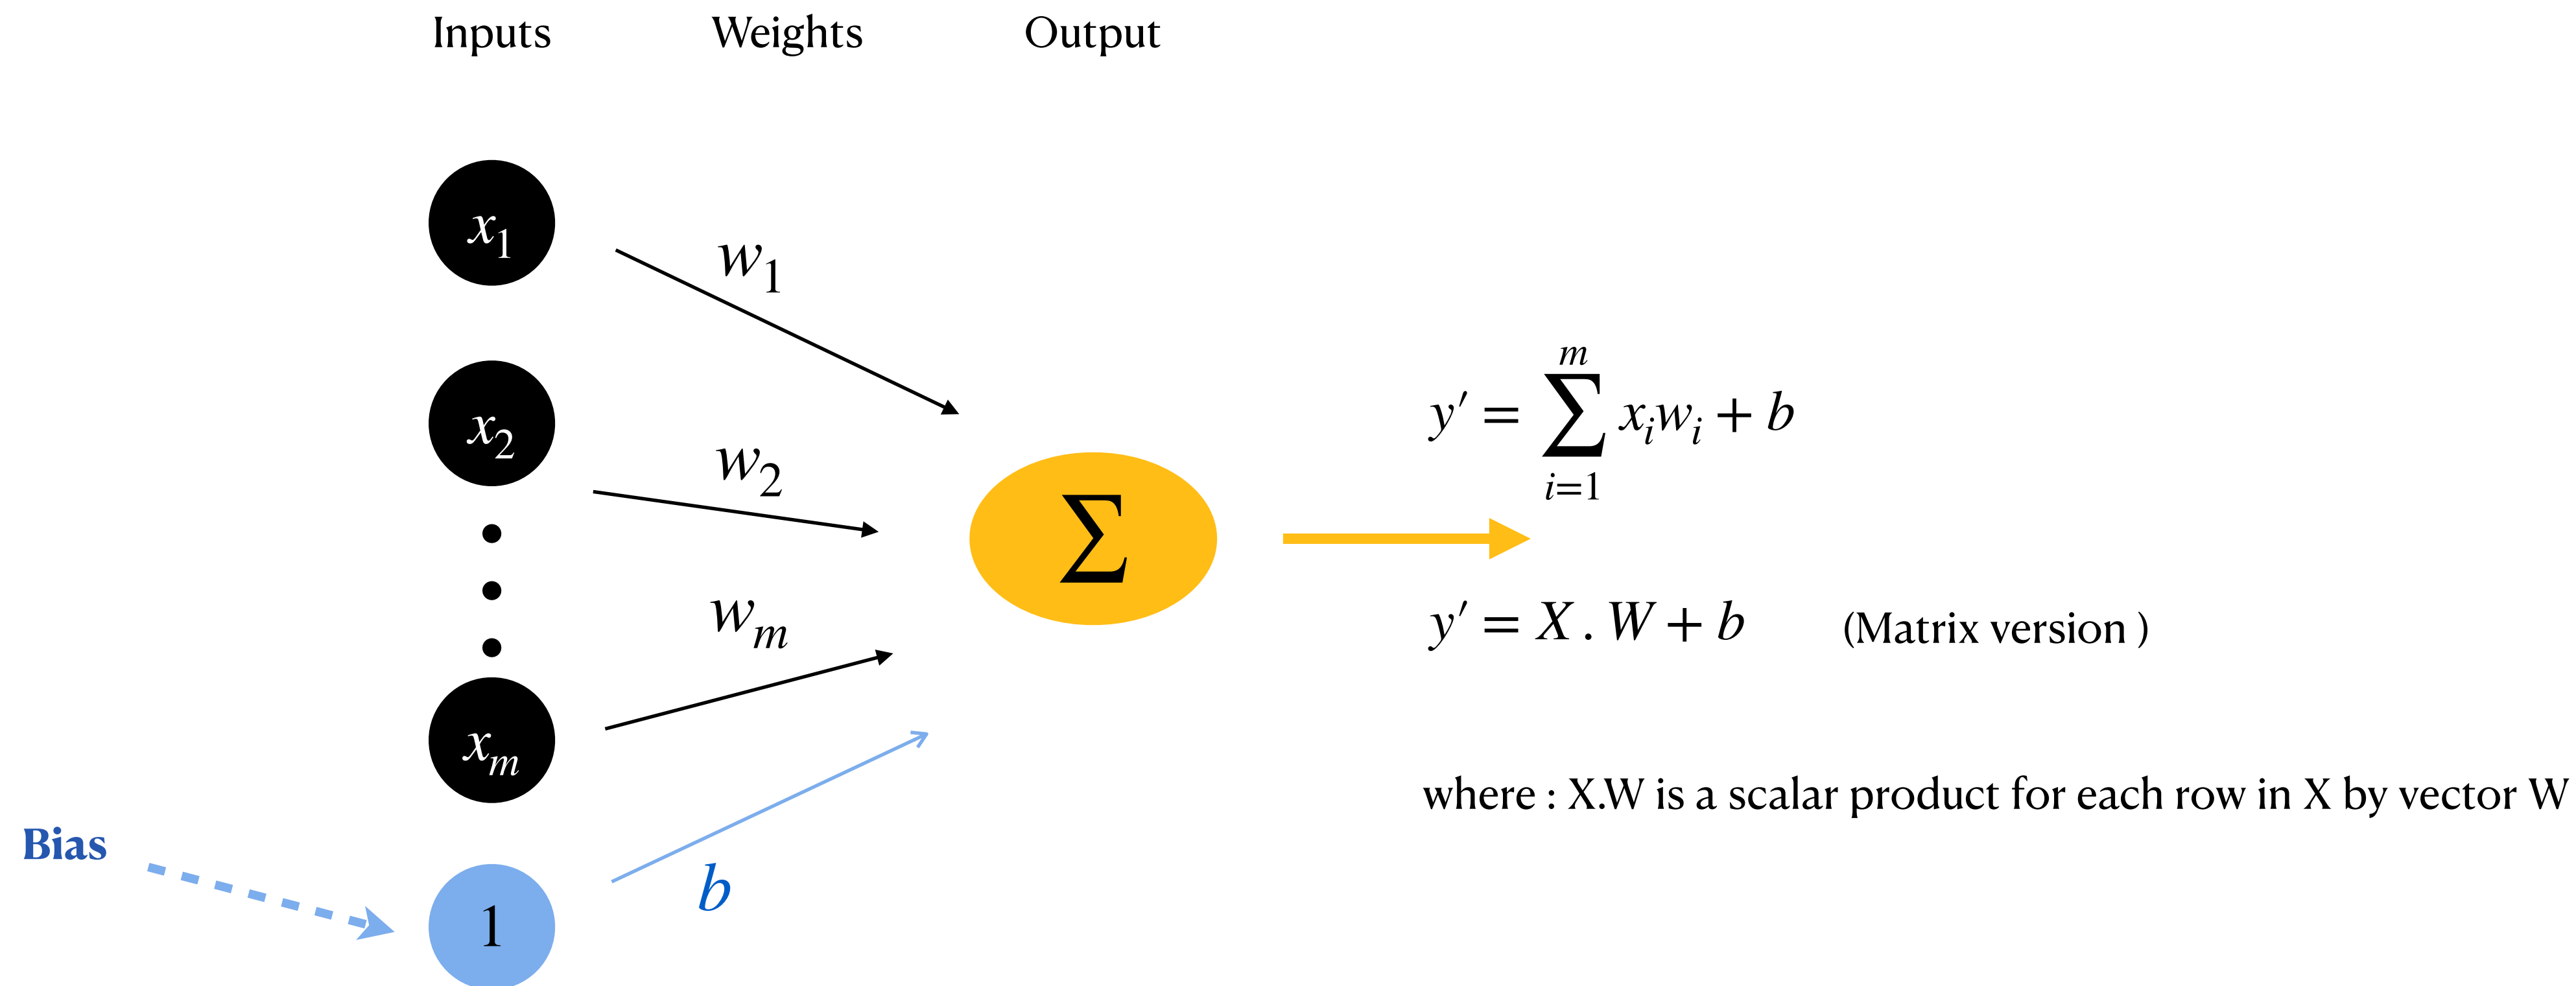
Formalism

- One of fundamental methods in *supervised Machine Learning*.
- Used to *predict* continuous values.
- Predict the value of a *dependent variable* based on the values of another *independent variables*.

Linear Regression

Formalism

- One of fundamental methods in **supervised Machine Learning**.
- Used to **predict** continuous values.
- Predict the value of a *dependent variable* based on the values of another *independent variables*.



Linear Regression

Formalism

- One of fundamental methods in *supervised Machine Learning*.
- Used to *predict* continuous values.
- Predict the value of a *dependent variable* based on the values of another *independent variables*.


★ General formulation, *the inference function* : $y = X \cdot W + b$

- **y** is the dependent variable : continuous values we want to predict
- **X** represent a data set of n samples and m features (independent variables).
- **W** is a vector of weights (*the unknown, we want to learn*), of size $(1,m)$.
- **b** is a scalar value called the bias (*unknown value, we want to learn*).

★ The goal is to find vector of weights **W** and scalar term **b**

Linear Regression


Formalism



Linear Regression

Formalism


- We want to find the equation fitting the *predicted line* to the *data*.
- Example of data with only one weight (i.e. one feature 'w')
- Linear regression consists in finding this *infinite red line* (a line not a circle, arc, or other shape ... otherwise we need non-linear methods).
- We call this line a *regression model*.



Linear Regression

Formalism


- Some value of weight 'w' and bias 'b'.
- These weights controls the shape of the line.



Linear Regression

Formalism


- Some value of weight 'w' and bias 'b'.
- These weights controls the shape of the line.



Linear Regression

Formalism


- Some value of weight ' w ' and bias ' b '.
- These weights controls the shape of the line.



Linear Regression

Formalism


- Some value of weight 'w' and bias 'b'.
- These weights controls the shape of the line.
- Almost !



Linear Regression

Formalism


- Not easy to find the exact position of the line !
- We use an optimisation algorithm known as "**Gradient Descent Algorithm**".
- We can observe that the most fitting line is truly "*close*" to all the data points, especially when compared to the initial line (slide 6).



Linear Regression

Formalism


- Not easy to find the exact position of the line !
- We use an optimisation algorithm known as "**Gradient Descent Algorithm**".
- We can observe that the most fitting line is truly "*close*" to all the data points, especially when compared to the initial line (slide 6).
- This closeness is mathematically quantified by *a loss (or error) function*.
- The objective of GDA algorithm is to *minimize* this specified loss function.



Linear Regression

Loss function

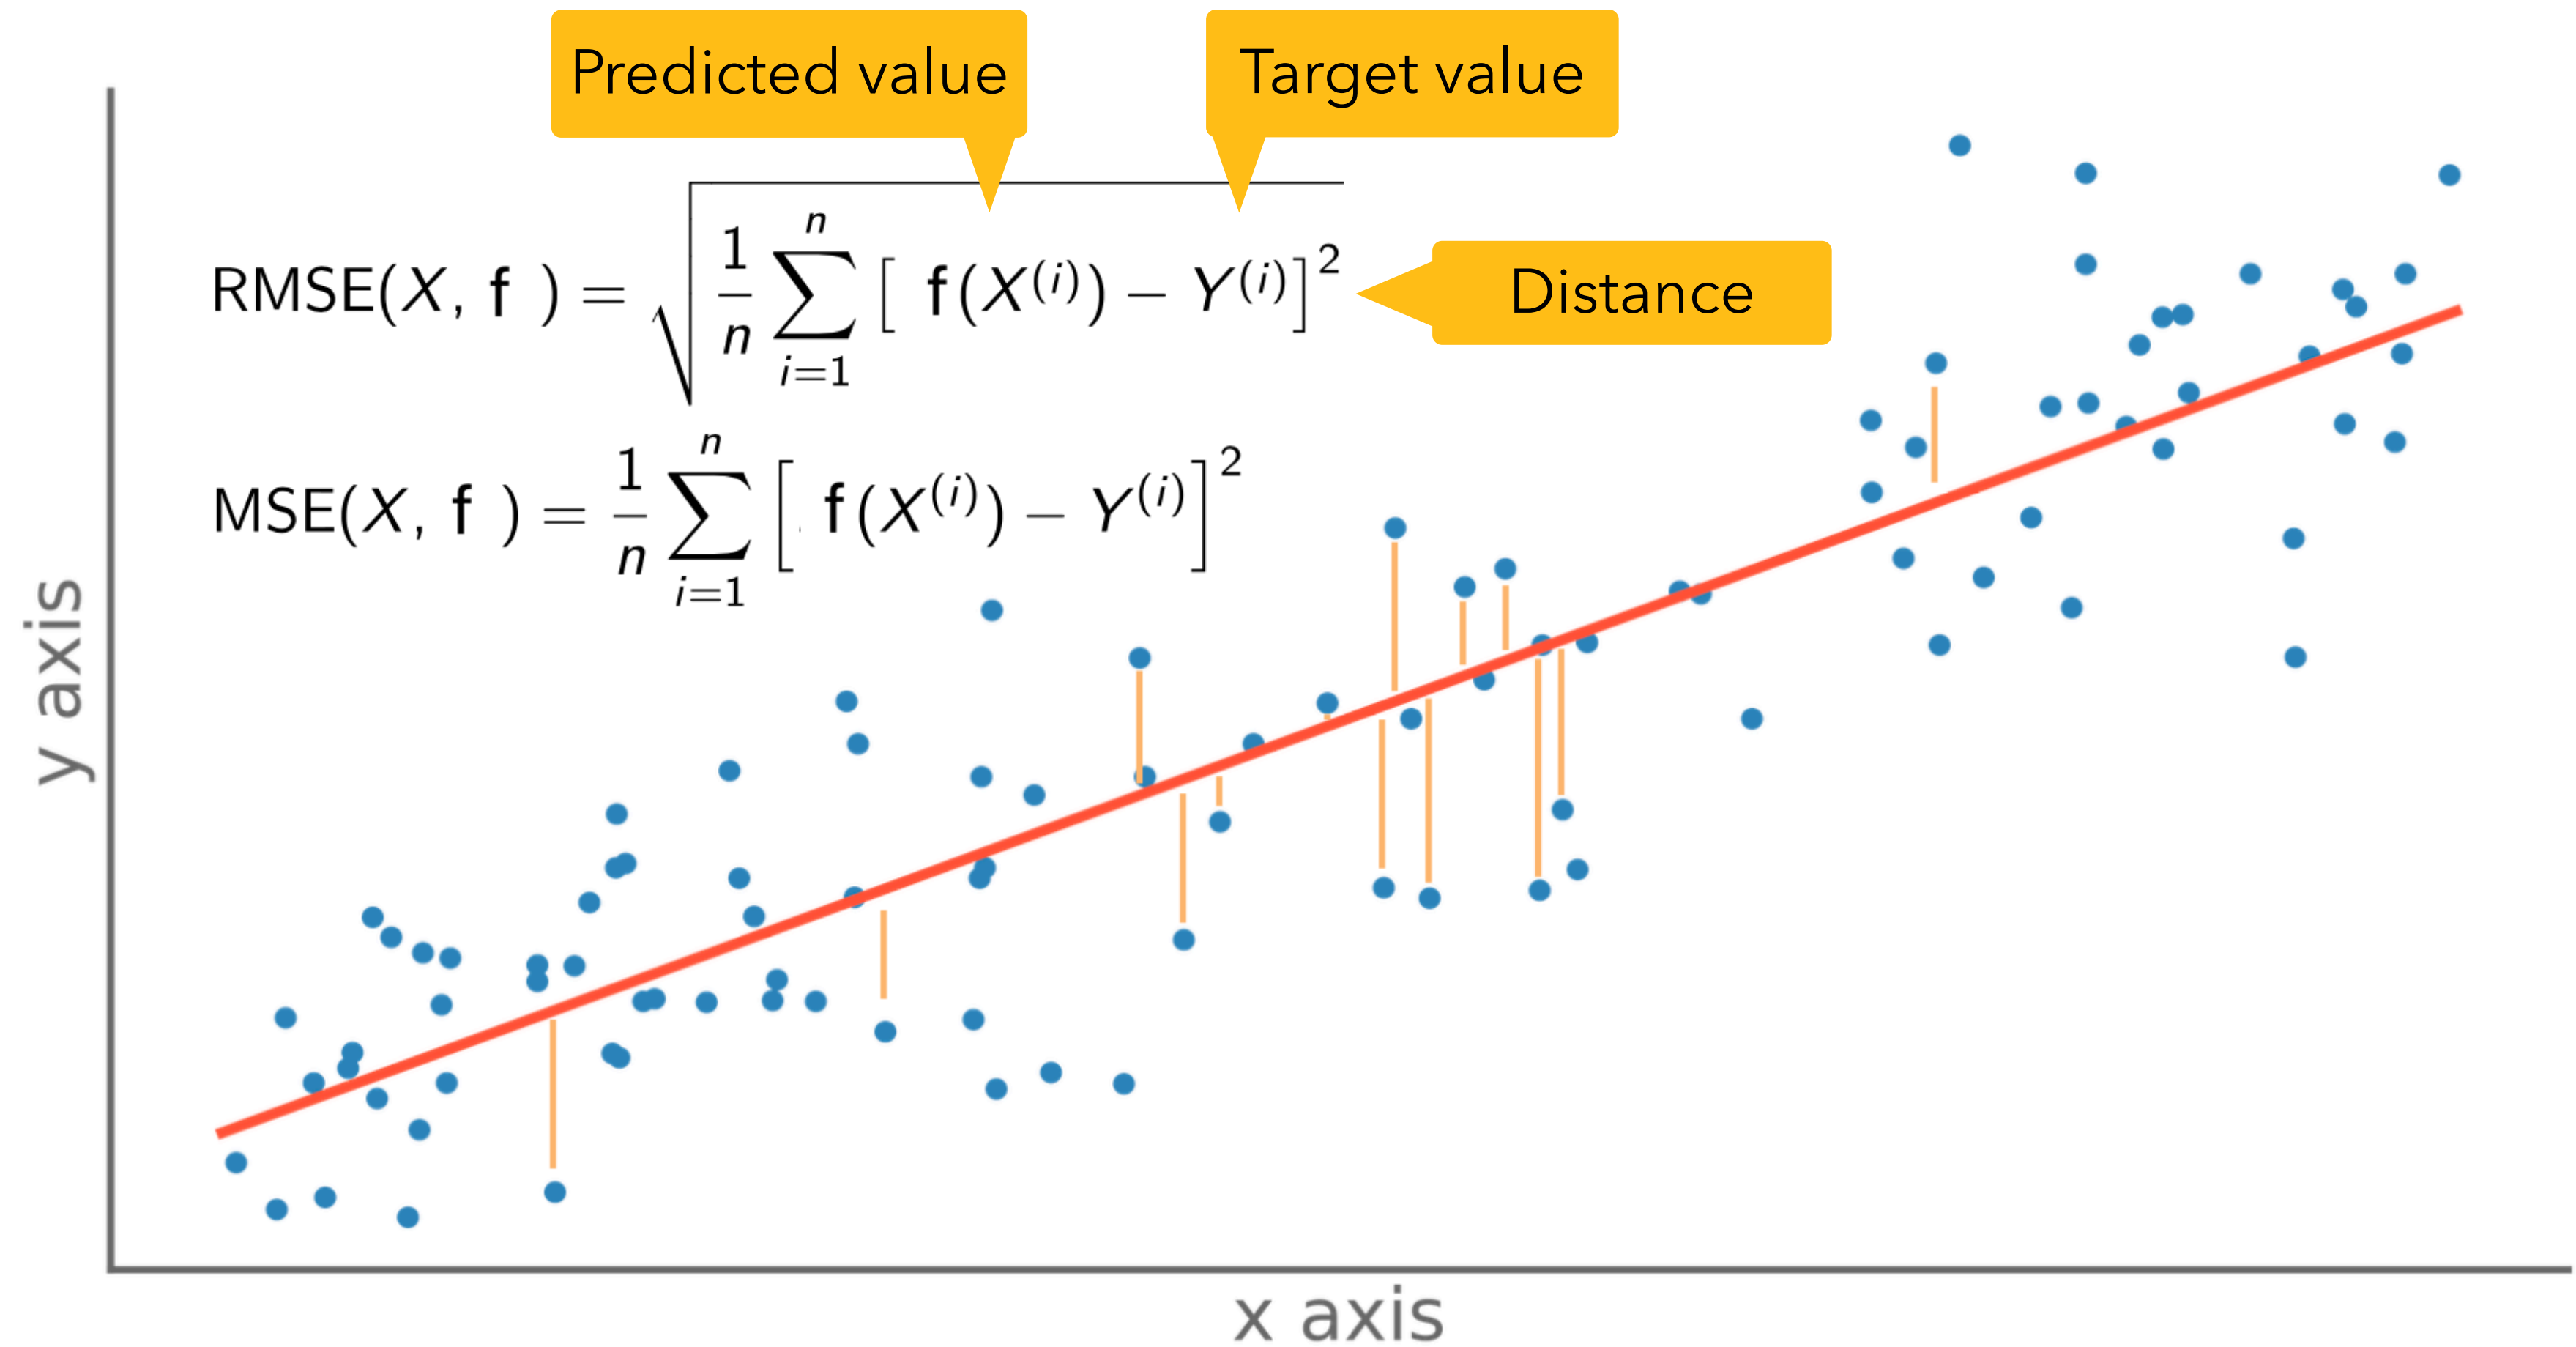
- Some of the most used error functions in LR :
 - * MSE : Mean Square Error function.
 - * RMSE : Root Mean Square Error function.



Linear Regression

Loss function


- Some of the most used error functions in LR :
- * MSE : Mean Square Error function.
- * RMSE : Root Mean Square Error function.



Linear Regression

Loss function


- Some value of weight 'w' and bias 'b'.
- These weights controls the shape of the line.
- $MSE = 4600$



Linear Regression

Loss function


- Some value of weight 'w' and bias 'b'.
- These weights controls the shape of the line.
- $MSE = 2495$ ↓



Linear Regression

Loss function


- Some value of weight 'w' and bias 'b'.
- These weights controls the shape of the line.
- $MSE = 1520$ ↓



Linear Regression

Loss function


- Some value of weight 'w' and bias 'b'.
- These weights controls the shape of the line.
- $MSE = 956$ ↓
- Almost !



Linear Regression

Loss function

- ***Optimal*** value of weight 'w' and bias 'b'.
- We can observe, that most fitted line is really ***"close"*** to the whole data points, compared to the first line (in slide 6).
- $MSE = 903$ ★



Linear Regression

Gradient Descent Algorithm

- In the previous slide, we observed that GDA algorithm progressively converge toward optimal weights, corresponding to the *minimum mean squared error (MSE)* value.
- Before introducing this algorithm, let's first discuss the underlying intuition.

Linear Regression

Gradient Descent Algorithm

- In the previous slide, we observed that GDA algorithm progressively converge toward optimal weights, corresponding to the **minimum mean squared error (MSE)** value.
- Before introducing this algorithm, let's first discuss the underlying intuition.

- ▶ We first define $J(W) = \frac{1}{n} \sum_{i=1}^n L(y_i, y'_i) = \frac{1}{n} \sum_{i=1}^n [y_i - y'_i]^2$ (i.e. MSE equation)
- ▶ In matrix form : $J(W) = \frac{1}{n} \sum_{i=1}^n [Y - (X \cdot W + b)]^2$
- ▶ We want to find weights W (and bias b) that **minimizes** $J(W)$, how ?
- ▶ Randomly ... but we can find a better and direct way : **gradient calculation**

Linear Regression

Gradient Descent Algorithm

- In the previous slide, we observed that GDA algorithm progressively converge toward optimal weights, corresponding to the **minimum mean squared error (MSE)** value.
- Before introducing this algorithm, let's first discuss the underlying intuition.

- ▶ We first define $J(W) = \frac{1}{n} \sum_{i=1}^n L(y_i, y'_i) = \frac{1}{n} \sum_{i=1}^n [y_i - y'_i]^2$ (i.e. MSE equation)
- ▶ In matrix form : $J(W) = \frac{1}{n} \sum_{i=1}^n [Y - (X \cdot W + b)]^2$
- ▶ We want to find weights W (and bias b) that **minimizes** $J(W)$, how ?
- ▶ Randomly ... but we can find a better and direct way : **gradient calculation**

- The gradient of a function at a given point shows the direction of the fastest increase in the function's value.
- Gradient descent uses this property to find the function's minimum by systematically moving to the **opposite direction**.

Linear Regression

Gradient Descent Algorithm

- In the previous slide, we observed that GDA algorithm progressively converge toward optimal weights, corresponding to the **minimum mean squared error (MSE)** value.
- Before introducing this algorithm, let's first discuss the underlying intuition.

► Gradient of **weights** :

$$\begin{aligned}\frac{\delta J(W)}{\delta W} &= \frac{\delta}{\delta W} \left[\frac{1}{n} \sum_{i=1}^n [y_i - (x_i \cdot W + b)]^2 \right] \\ &= \frac{2}{n} \sum_{i=1}^n -x_i [y_i - (x_i \cdot W + b)] \\ &= \frac{2}{n} \sum_{i=1}^n -x_i [y_i - y'_i] \\ &= -\frac{2}{n} X^T (Y - Y') \quad (\text{matrix form})\end{aligned}$$

► Gradient of **bias** :

$$\begin{aligned}\frac{\delta J(W)}{\delta b} &= \frac{\delta}{\delta b} \left[\frac{1}{n} \sum_{i=1}^n [y_i - (x_i \cdot W + b)]^2 \right] \\ &= \frac{2}{n} \sum_{i=1}^n [y_i - (x_i \cdot W + b)] \\ &= \frac{2}{n} \sum_{i=1}^n -[y_i - y'_i] \\ &= -\frac{2}{n} (Y - Y') \quad (\text{matrix form})\end{aligned}$$

Linear Regression

Gradient Descent Algorithm

- ➡ Entry : Data input values X of size (n samples, m features), data target values Y of size $(n, 1)$
MaxIterations, learning rate α
- ➡ Output : Weights W of size $(1, m)$, and bias b

❖ Training (Learning weights)

(1) **Initialisation** : random values for vector W , and bias b

(2) **While** termination condition :

For each sample (x_i, y_i) :

- Calculate predicted value $y'_i = x_i \cdot W + b$
- Compute gradient $\frac{\delta J(W)}{\delta W}$ and $\frac{\delta J(W)}{\delta b}$
- Updating weights : $W = W - \alpha \frac{\delta J(W)}{\delta W}$
- Updating bias : $b = b - \alpha \frac{\delta J(W)}{\delta b}$

(3) **Return** W, b

Linear Regression

Gradient Descent Algorithm

- ➡ Entry : Data input values X of size (n samples, m features), data target values Y of size $(n, 1)$
MaxIterations, learning rate α
- ➡ Output : Weights W of size $(1, m)$, and bias b

❖ Training (Learning weights)

- (1) **Initialisation** : random values for vector W , and bias b
- (2) **While** termination condition :

For each sample (x_i, y_i) :

- Calculate predicted value $y'_i = x_i \cdot W + b$
- Compute gradient $\frac{\delta J(W)}{\delta W}$ and $\frac{\delta J(W)}{\delta b}$
- Updating weights : $W = W - \alpha \frac{\delta J(W)}{\delta W}$
- Updating bias : $b = b - \alpha \frac{\delta J(W)}{\delta b}$

- (3) **Return** W, b

❖ Learning rate $\alpha \in [0,1]$

- needs to be adjusted carefully
- Big value => Big updating steps
- Small value => Small updating steps

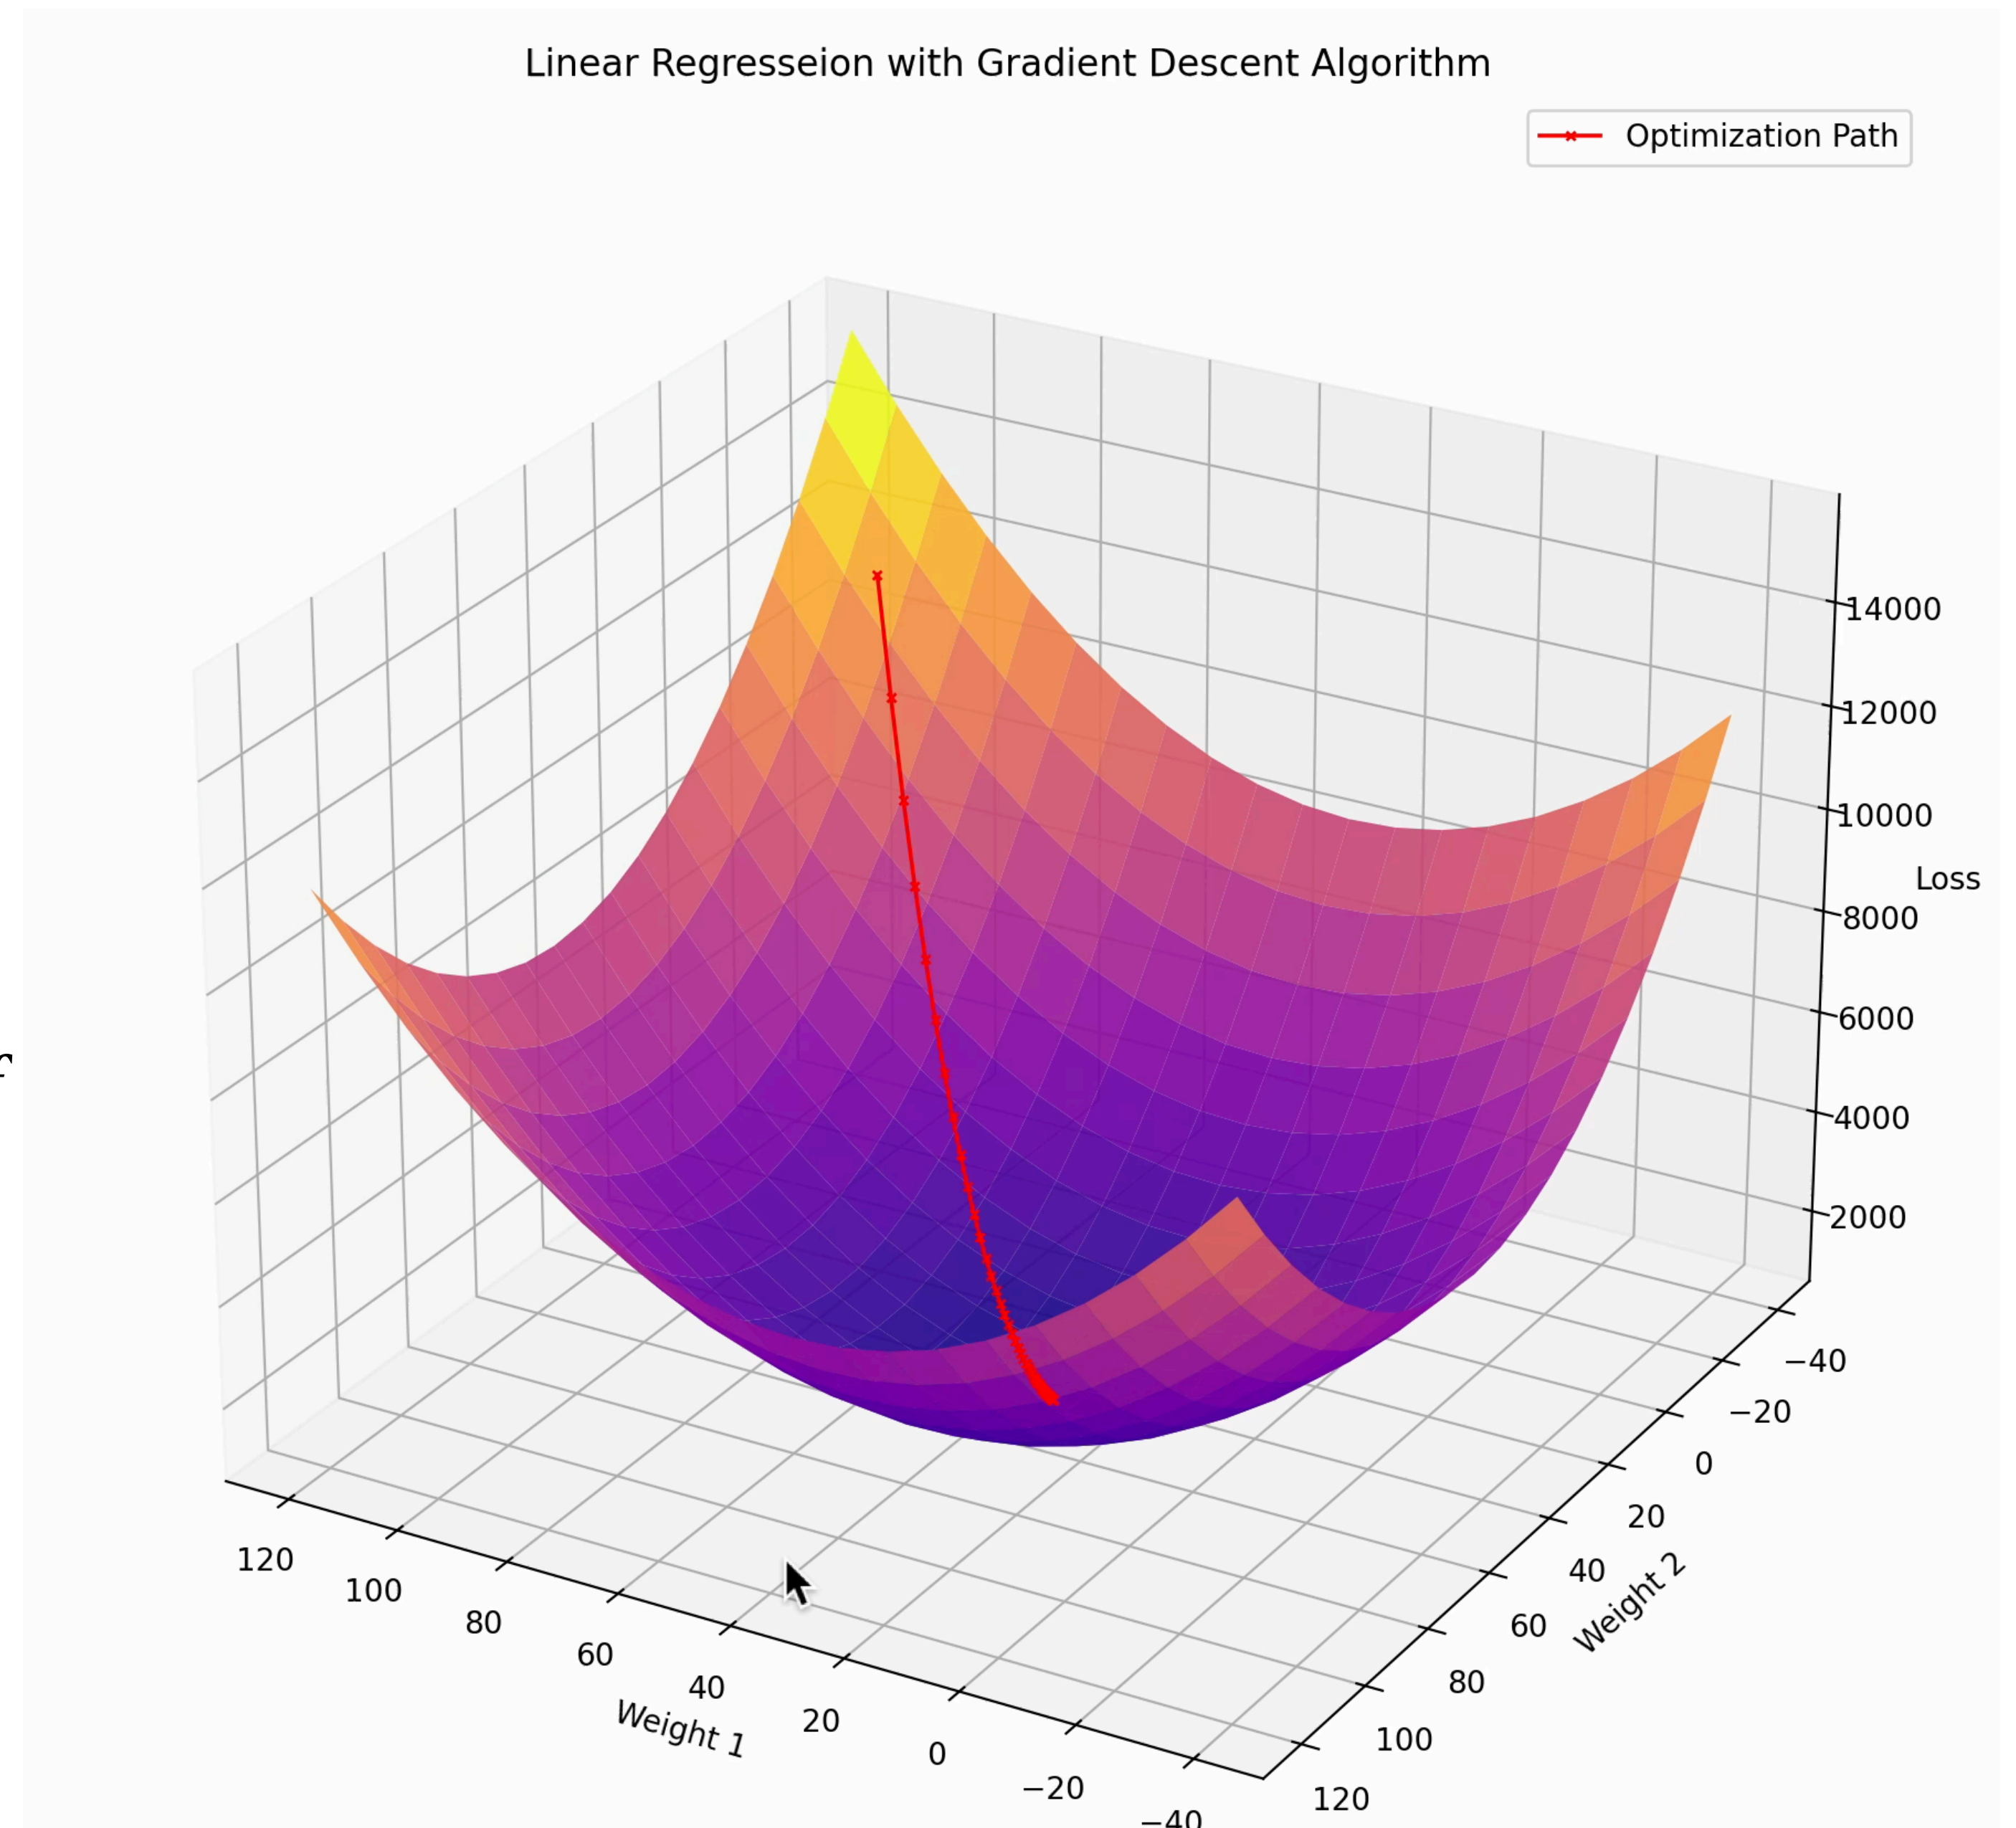
★ Prediction (Testing model)

Simply calculate $y'_i = x_i \cdot W + b$

Linear Regression

GDA & Loss landscape

- Observe here, a model with **two weights** (not a single weight as in the former example!).
- For each pair (weight 1, weight 2), MSE value (third axis) is calculated. Obtaining **MSE landscape**.
- So the **GDA** algorithm : starts from some value of weights, then minimize progressively MSE error until arriving to optimal value.
- The GDA algorithm trajectory is plotted in red.
- **MSE** is a convex function !



Practical Activity 1 (part I)

Linear Regression

- Download **LAB1** support from Moodle : <https://moodle.myefrei.fr/course/view.php?id=14646>
- LABs comprise groups of up to **2** students.
- Bonuses can be attributed ! if the teacher considers the results to be relevant.

Logistic Regression

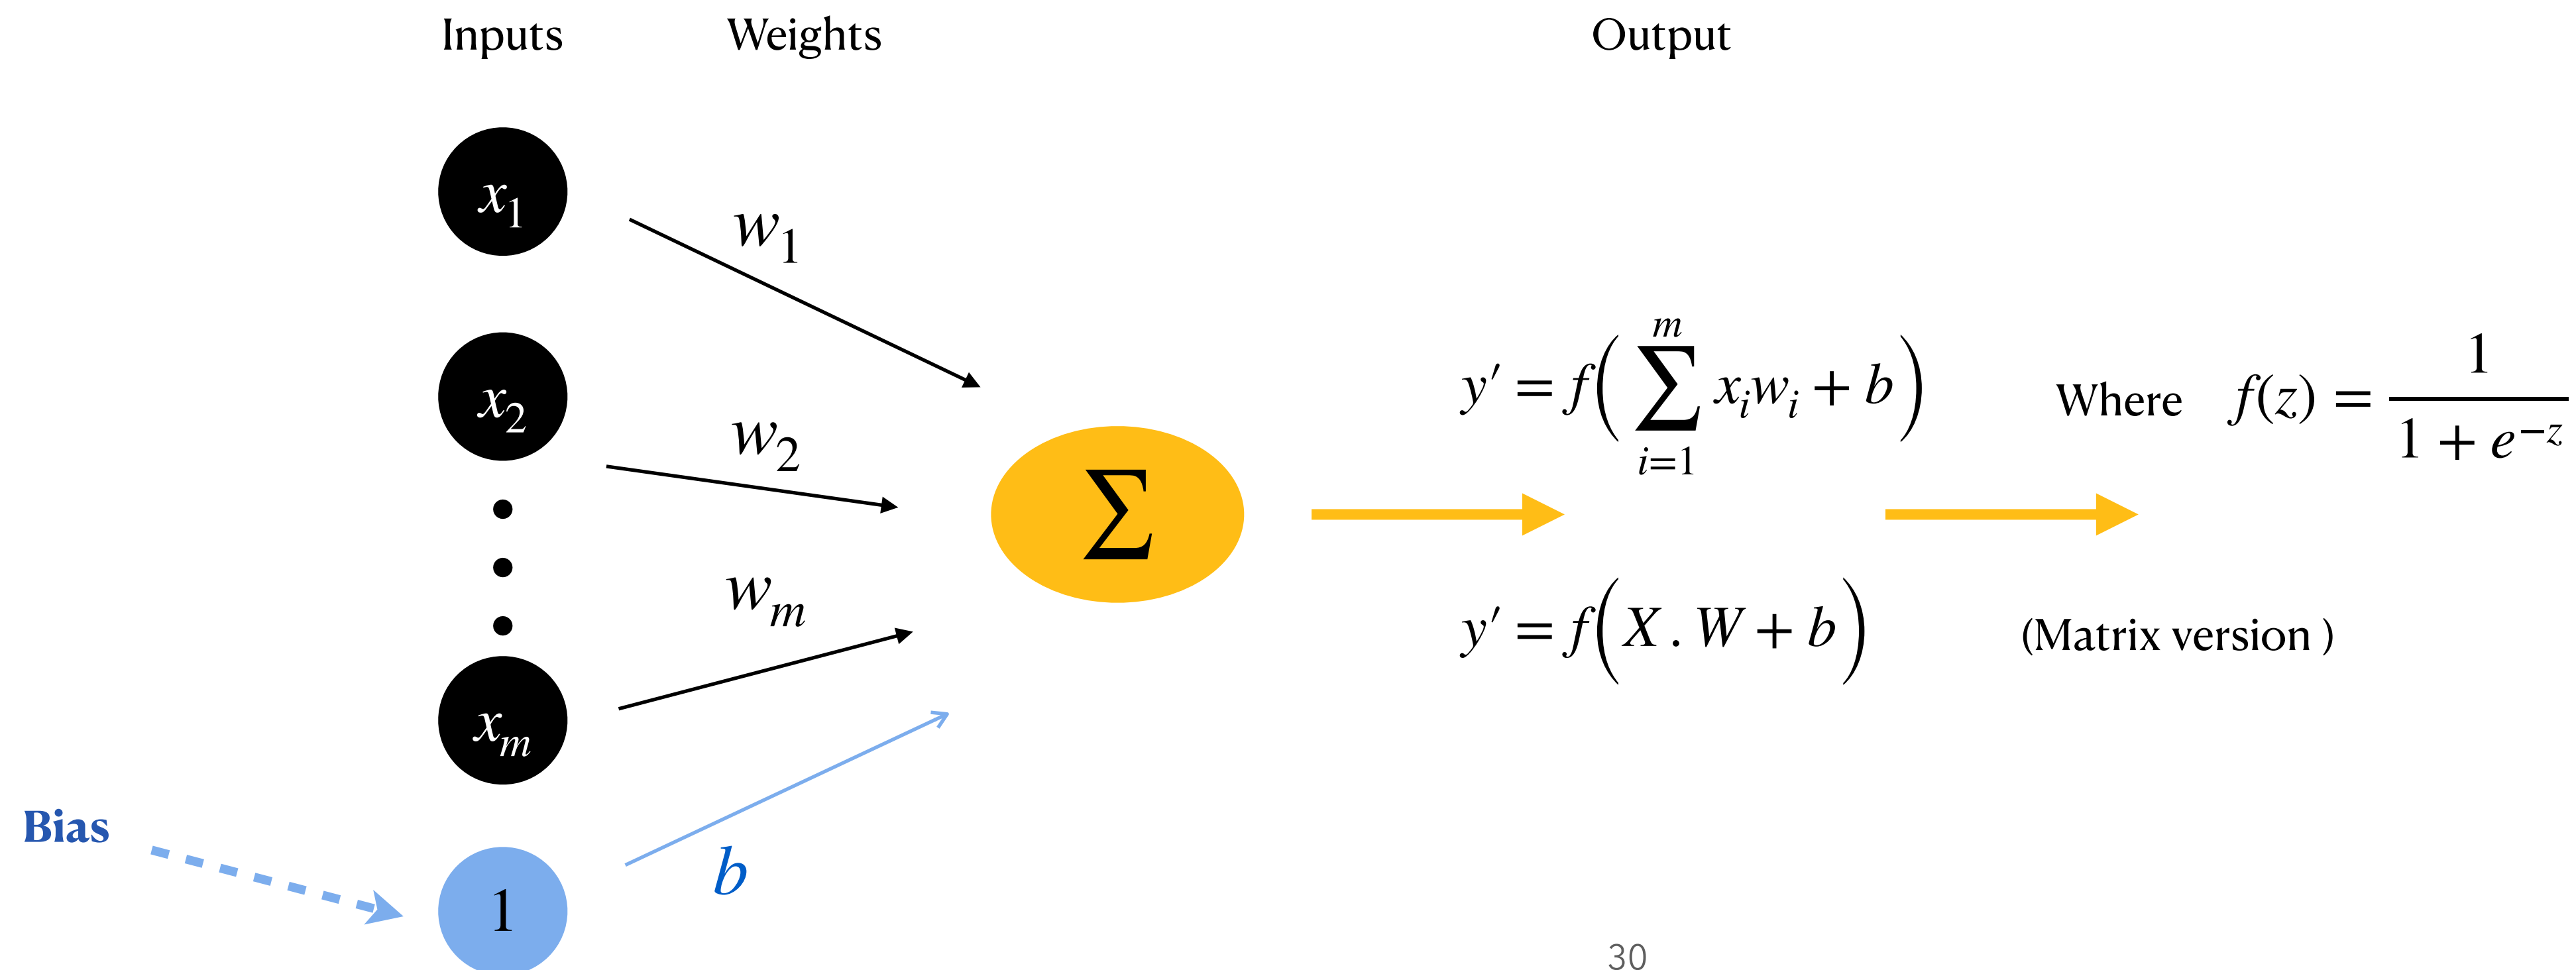
Formalism

- One of fundamental methods in *supervised Machine Learning*.
- Used for *binary classification tasks* !
- Outcome variable is categorical with two possible outcomes (e.g., 0 or 1)
- It models the probability that an observation belongs to one of the categories.

Logistic Regression

Formalism

- One of fundamental methods in *supervised Machine Learning*.
- Used for *binary classification tasks* !
- Outcome variable is categorical with two possible outcomes (e.g., 0 or 1)
- It models the probability that an observation belongs to one of the categories.



Logistic Regression

Formalism

- One of fundamental methods in *supervised Machine Learning*.
- Used for *binary classification tasks* !
- Outcome variable is categorical with two possible outcomes (e.g., 0 or 1)
- It models the probability that an observation belongs to one of the categories.

★ General formulation, *the inference function* : $y = f(X \cdot W + b)$ where $f(z) = \frac{1}{1 + e^{-z}}$

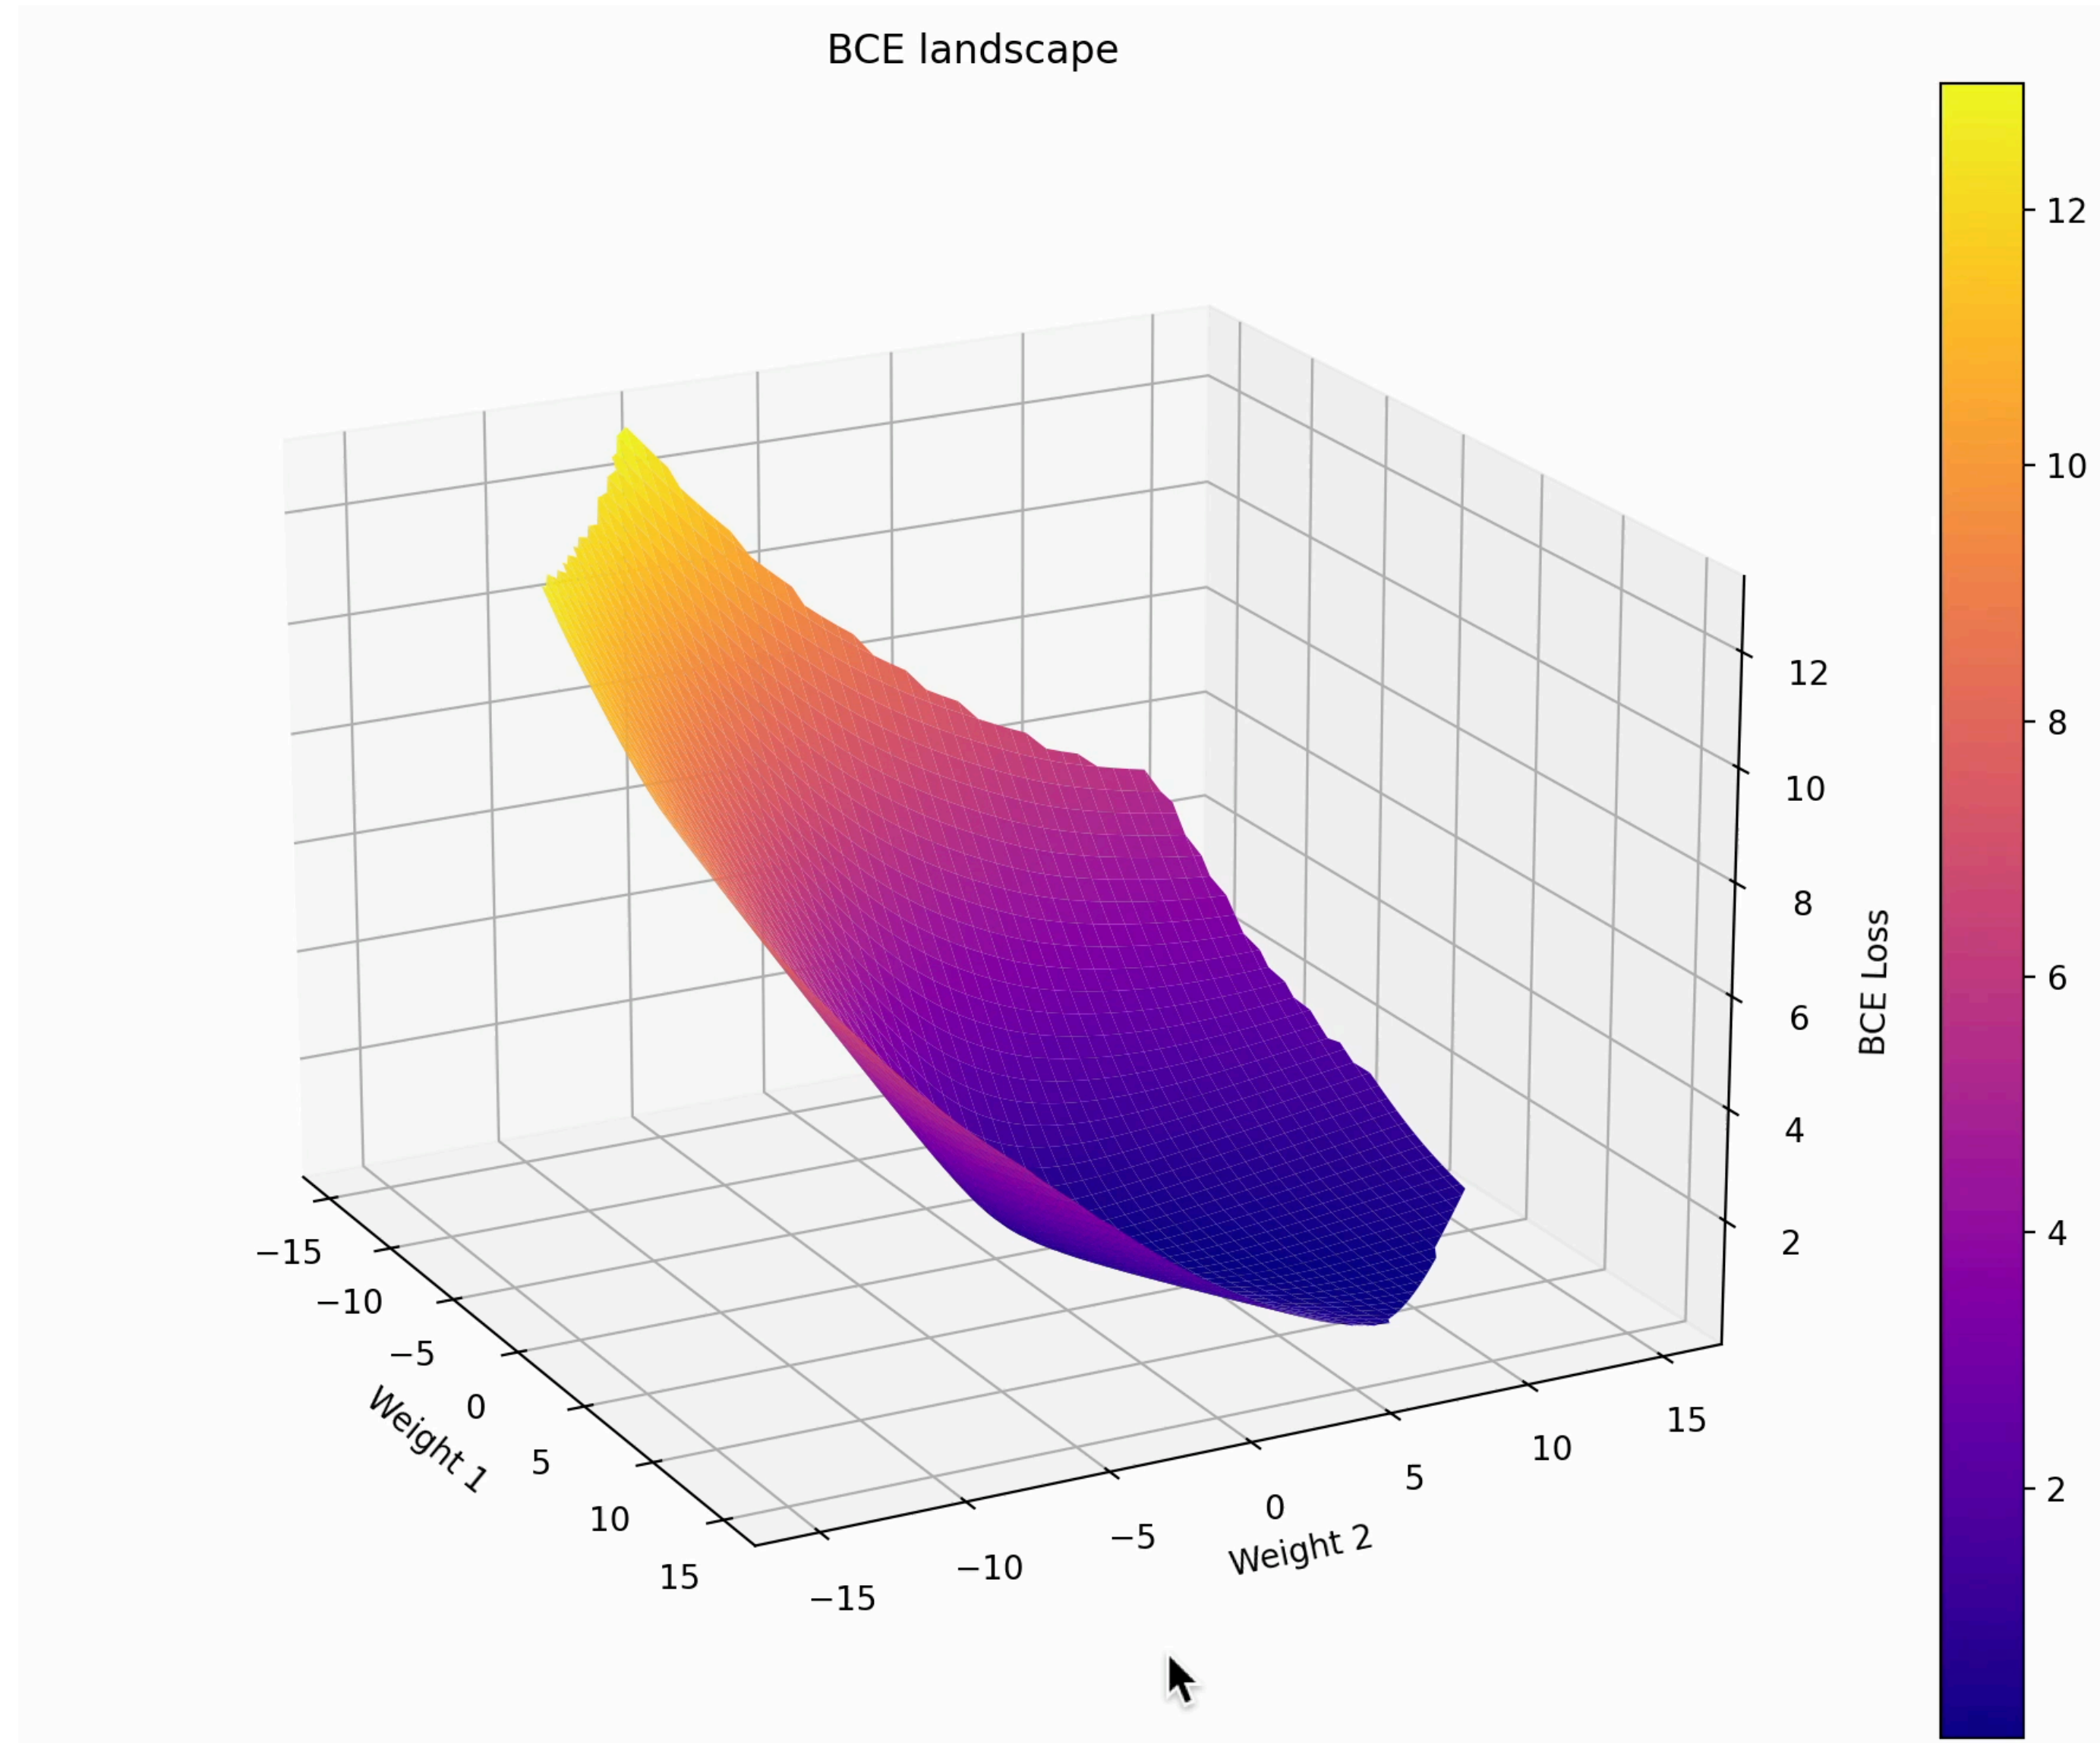
- **y** is the dependent variable : continuous values we want to predict
- **X** represent a data set of n samples and m features (independent variables).
- **W** is a vector of weights (*the unknown, we want to learn*), of size $(1, m)$.
- **b** is a scalar value called the bias (*unknown value, we want to learn*).

★ The goal is to find vector of weights **W** and scalar term **b**

Logistic Regression

Loss function

- Logistic regression does not use MSE because its output is probabilistic bounded between 0 and 1.
- It uses a loss function known as the log loss (or **Binary Cross-Entropy loss**) for binary classification.
- It measures the penalty for a given class prediction in terms of the distance from the actual label.
- ➔ The formula :
$$J(W) = -\frac{1}{n} \sum_{i=1}^n \left[y_i \log(y'_i) + (1 - y_i) \log(1 - y'_i) \right]$$
- ➔ **BCE** is a convex function, for simple LR case !



Logistic Regression

Gradient Descent Algorithm

- As for *MSE* error, *BCE* is also convex which ensures that Gradient Descent Algorithm has a global minimum
- Before introducing this algorithm, let's first discuss the underlying intuition.

- ▶ We first define : $J(W) = -\frac{1}{n} \sum_{i=1}^n \left[y_i \log(y'_i) + (1 - y_i) \log(1 - y'_i) \right]$ (i.e. BCE equation)
- ▶ Where : $y'_i = \frac{1}{1 + e^{-(X.W+b)}}$ (i.e. the sigmoid function)
- ▶ We want to find weights W (and bias b) that **minimizes** $J(W)$, we calculate gradients :
- ▶ Gradient of weights : $\frac{\delta J(W)}{\delta W} = -\frac{1}{n} X^T (Y - Y')$
- ▶ Gradient of bias : $\frac{\delta J(W)}{\delta b} = -\frac{1}{n} (Y - Y')$

- The gradient of a function at a given point shows the direction of the fastest increase in the function's value.
- Gradient descent uses this property to find the function's minimum by systematically moving to the **opposite direction**.

Logistic Regression

Gradient Descent Algorithm

- ➡ Entry : Data input values X of size (n samples, m features), data target values Y of size $(n, 1)$
MaxIterations, learning rate α , *Threshold* $\in [0,1]$
- ➡ Output : Weights W of size $(1, m)$, and bias b

❖ Training (Learning weights)

(1) **Initialisation** : random values for vector W , and bias b

(2) **While** termination condition :

For each random sample (x_i, y_i) :

- Calculate predicted value $y'_i = \left[1 + e^{-(x_i \cdot W + b)} \right]^{-1}$
- Compute gradient $\frac{\delta J(W)}{\delta W}$ and $\frac{\delta J(W)}{\delta b}$
- Updating weights : $W = W - \alpha \frac{\delta J(W)}{\delta W}$
- Updating bias : $b = b - \alpha \frac{\delta J(W)}{\delta b}$

(3) **Return** W, b

Logistic Regression

Gradient Descent Algorithm

- ➡ Entry : Data input values X of size (n samples, m features), data target values Y of size $(n, 1)$
MaxIterations, learning rate α , *Threshold* $\in [0,1]$
- ➡ Output : Weights W of size $(1, m)$, and bias b

❖ Training (Learning weights)

- (1) **Initialisation** : random values for vector W , and bias b
- (2) **While** termination condition :

For each random sample (x_i, y_i) :

- Calculate predicted value $y'_i = \left[1 + e^{-(x_i \cdot W + b)} \right]^{-1}$
- Compute gradient $\frac{\delta J(W)}{\delta W}$ and $\frac{\delta J(W)}{\delta b}$
- Updating weights : $W = W - \alpha \frac{\delta J(W)}{\delta W}$
- Updating bias : $b = b - \alpha \frac{\delta J(W)}{\delta b}$

- (3) **Return** W, b

❖ Learning rate $\alpha \in [0,1]$

- needs to be adjusted carefully
- Big value => Big updating steps
- Small value => Small updating steps

★ Prediction (Testing model)

- Simply calculate $y'_i = \left[1 + e^{-(x_i \cdot W + b)} \right]^{-1} \in [0,1]$
- **if** $(y'_i \geq \textit{Threshold})$ **then** x_i **is a cat**
else x_i **is a dog**

★ Accuracy

An accuracy function may be used to verify the accuracy of the model : *predicted classes Vs target classes for each x_i .*

Summary : Distinctions

	Linear Regression	Logistic Regression
The purpose	Prediction of continuous values	Binary Classification
The model form	$X \cdot W + b$	$\left[1 + e^{-(X \cdot W + b)}\right]^{-1}$
The loss function	Mean Squared Error (MSE)	Binary Cross-Entropy (BCE)
The algorithm	GDA	GDA
Limitation	Data landscape must be linear	Linear separation of data

For more complicated data landscapes => Polynomial Regression or use of Neural Networks.

Summary : Use cases

Linear Regression	Logistic Regression
Prediction of <i>adult characteristics</i> based on parent's characteristics	<i>Fraud detection</i> : can help teams identify anomalies in data that may predict fraud
Predicting <i>product sales</i> volume as a function of price, time of year and store location	<i>Disease prediction</i> : Predicting whether a person will develop heart disease based on IMC, smoking habits and genetic predisposition
Predict the <i>price of an airline ticket</i> as a function of origin, destination, time of year and airline company	<i>Attrition prediction</i> : it could be useful for HR and management to know whether high-performing employees are at risk of leaving the organization

Practical Activity 1 (part II)

Logistic Regression

- Download **LAB1** support from Moodle : <https://moodle.myefrei.fr/course/view.php?id=14646>
- LABs comprise groups of up to **2** students.
- Bonuses can be attributed ! if the teacher considers the results to be relevant.

